Attack Resilience Comparison
Dataset: UNSW | Clients: 11 | α=0.5 (Dirichlet) | Attack: grad_ascent | Seeds: n=5

Detection Performance vs Adversary Percentage (Macro-F1, 95% CI)

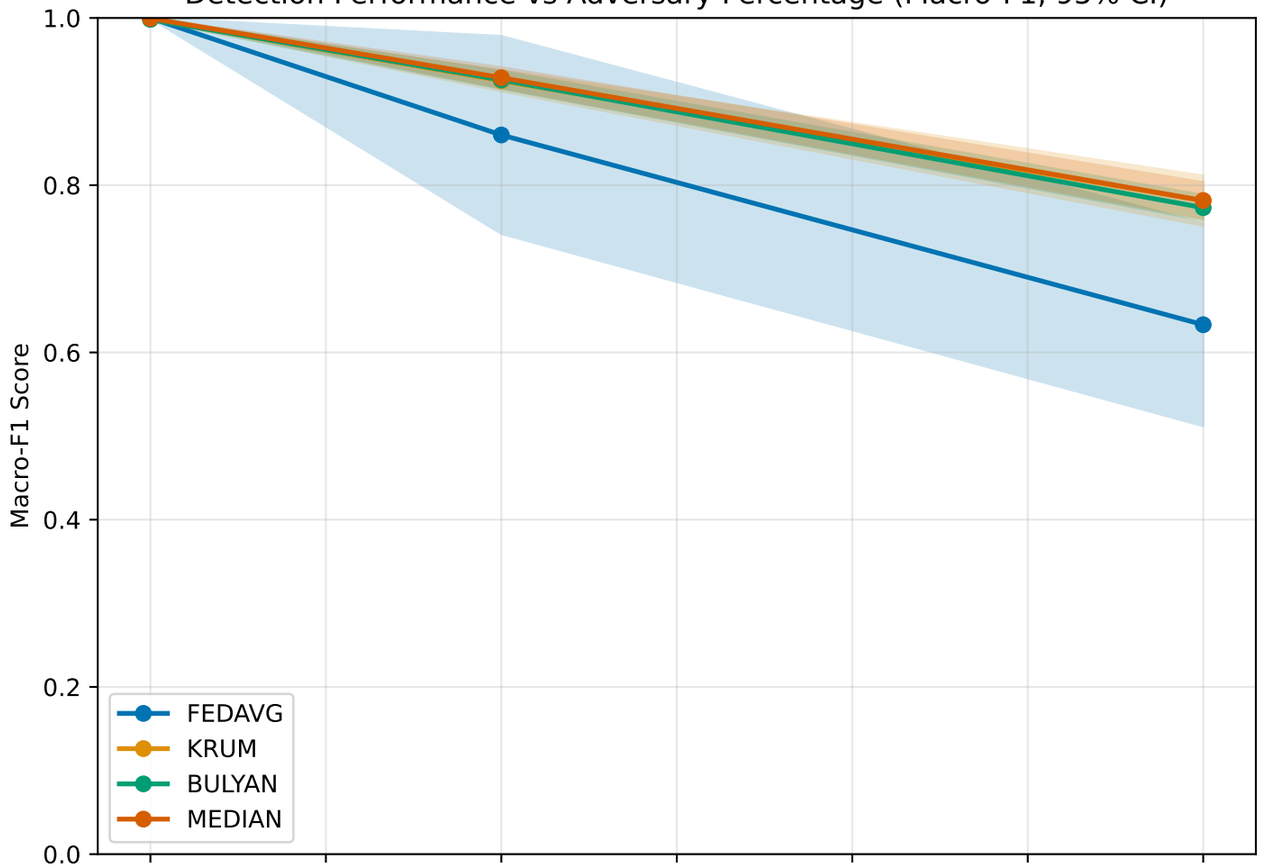

Supplementary: Model Drift (L2 Distance)

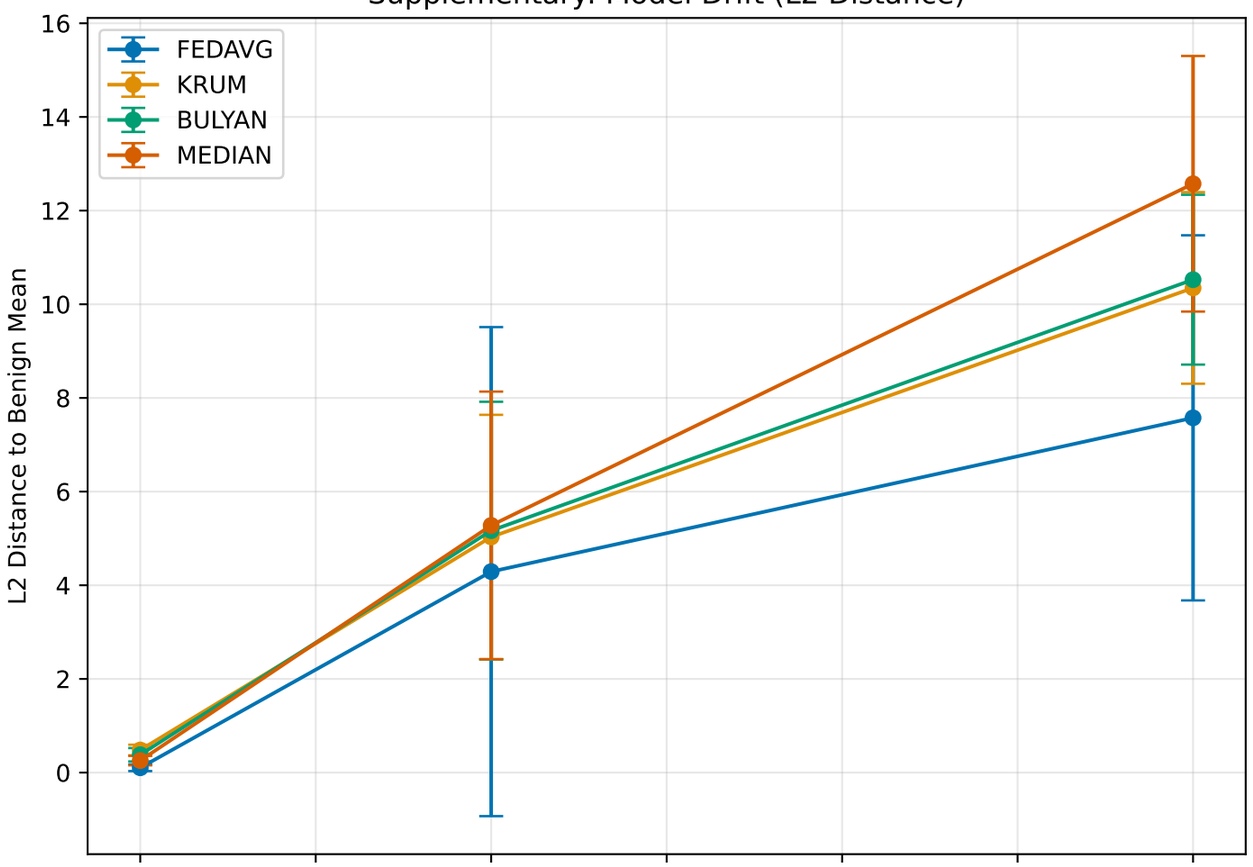

Performance Degradation Under Attack (95% CI, Bounded [0,100])

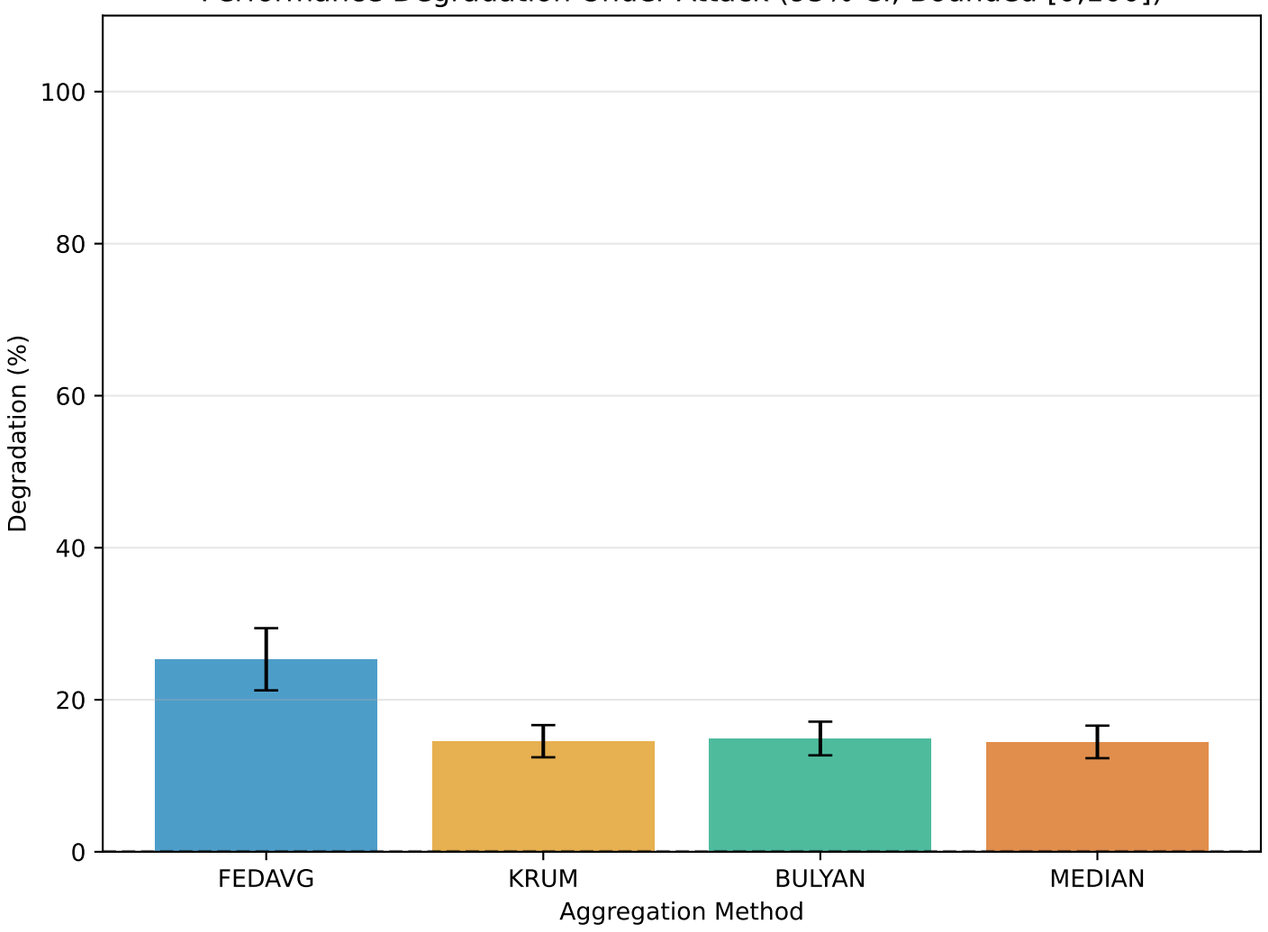

Supplementary: Model Alignment (Cosine Similarity)