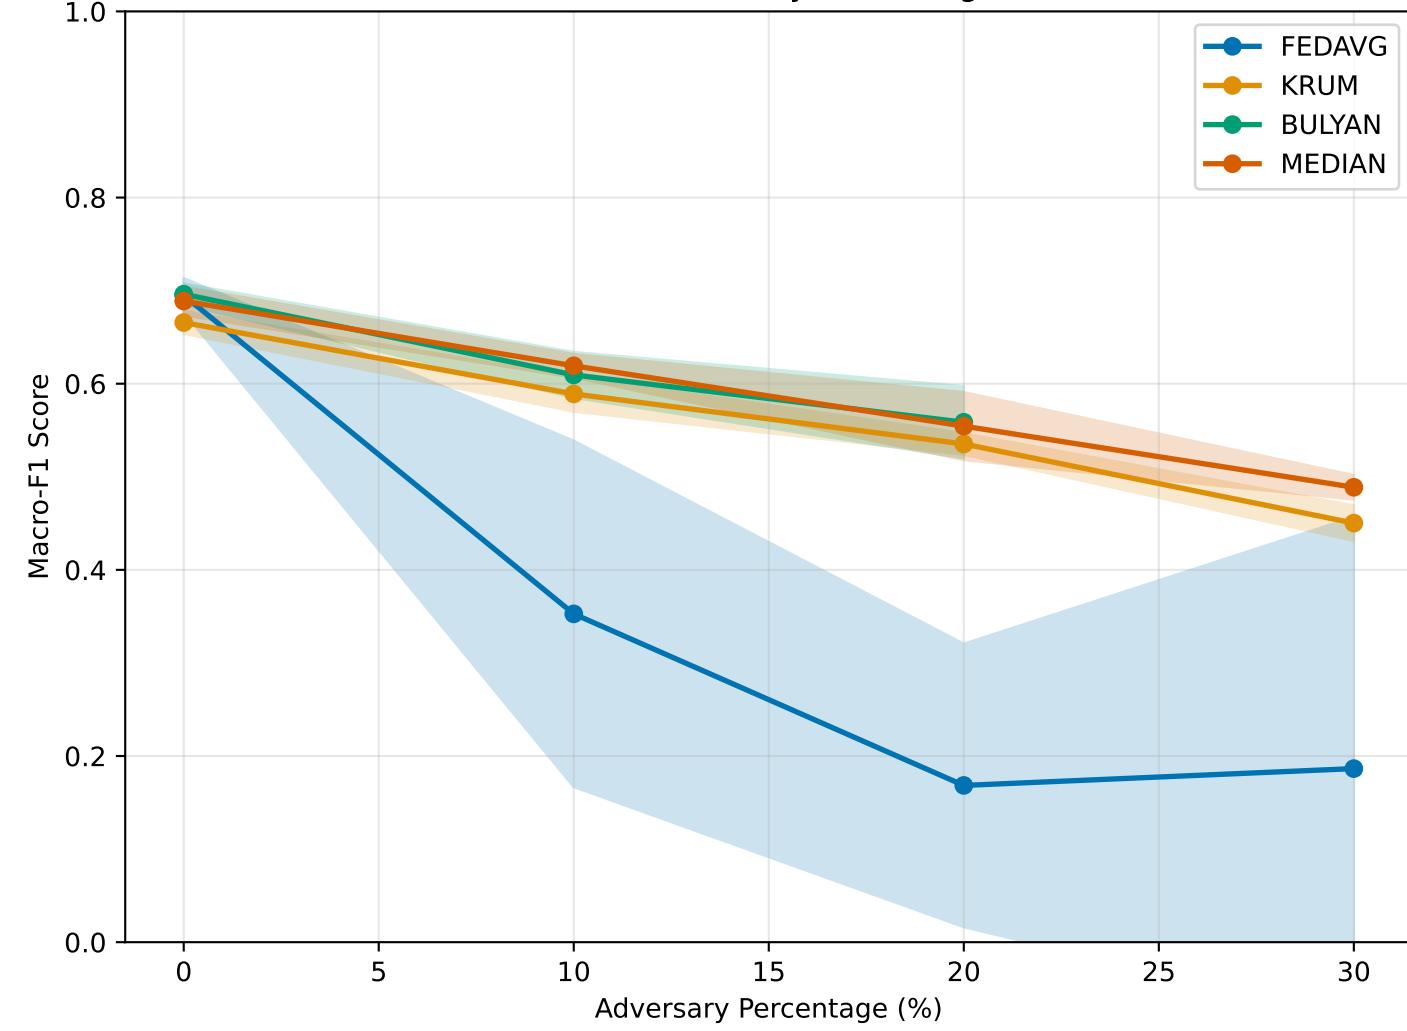


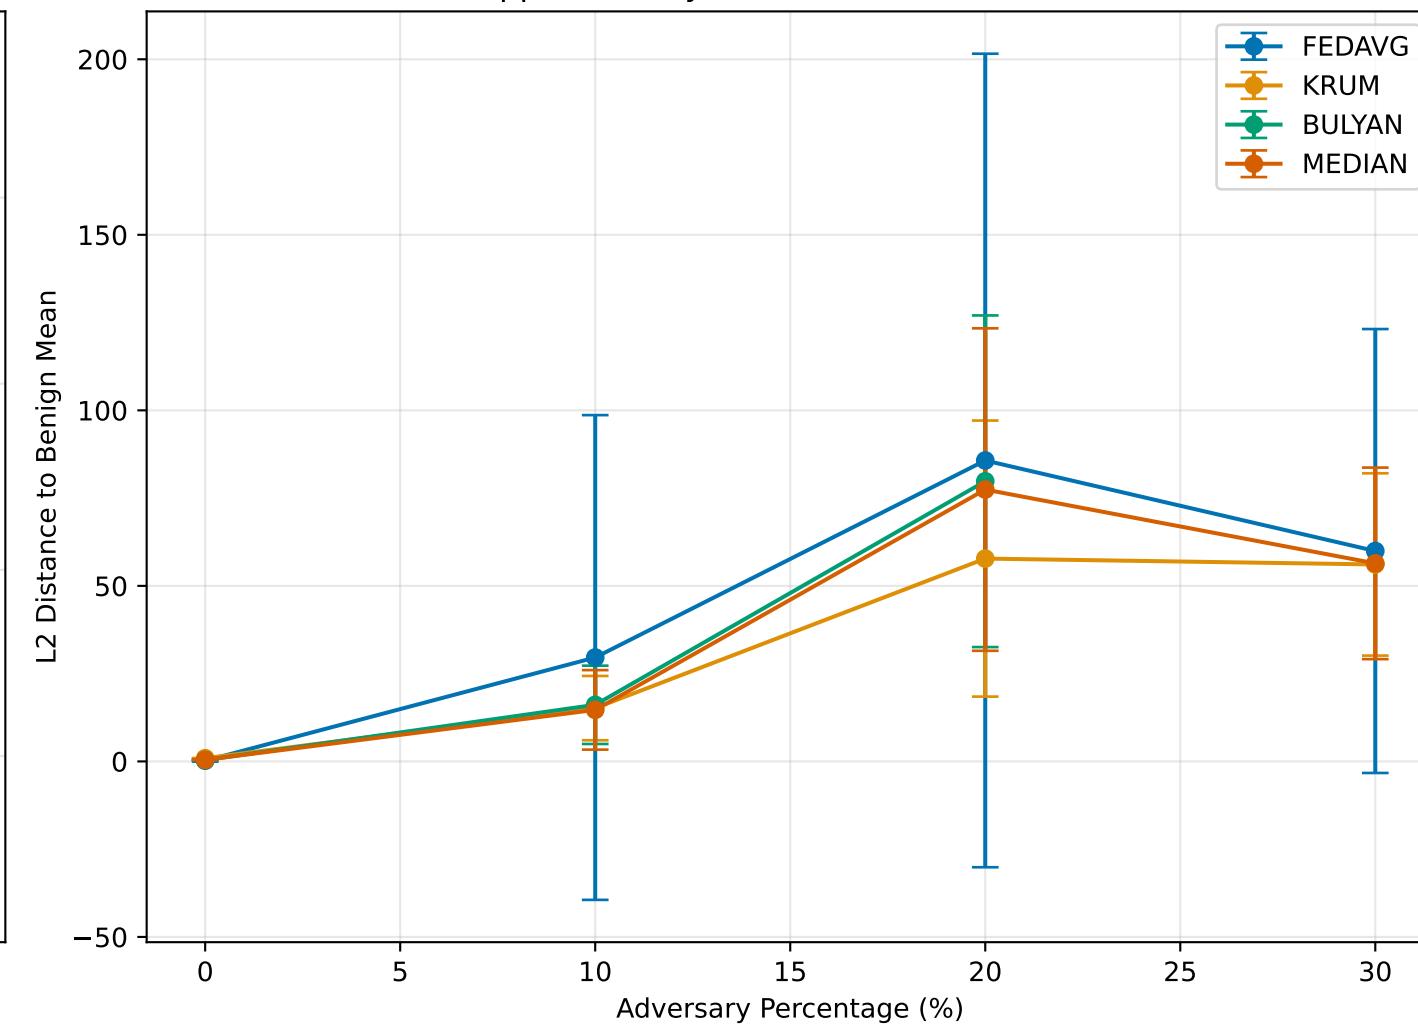
Attack Resilience Comparison

Dataset: EDGE-IIOTSET-QUICK | Clients: 3 | $\alpha=0.02$ (Dirichlet) | Attack: grad_ascent | Seeds: n=10

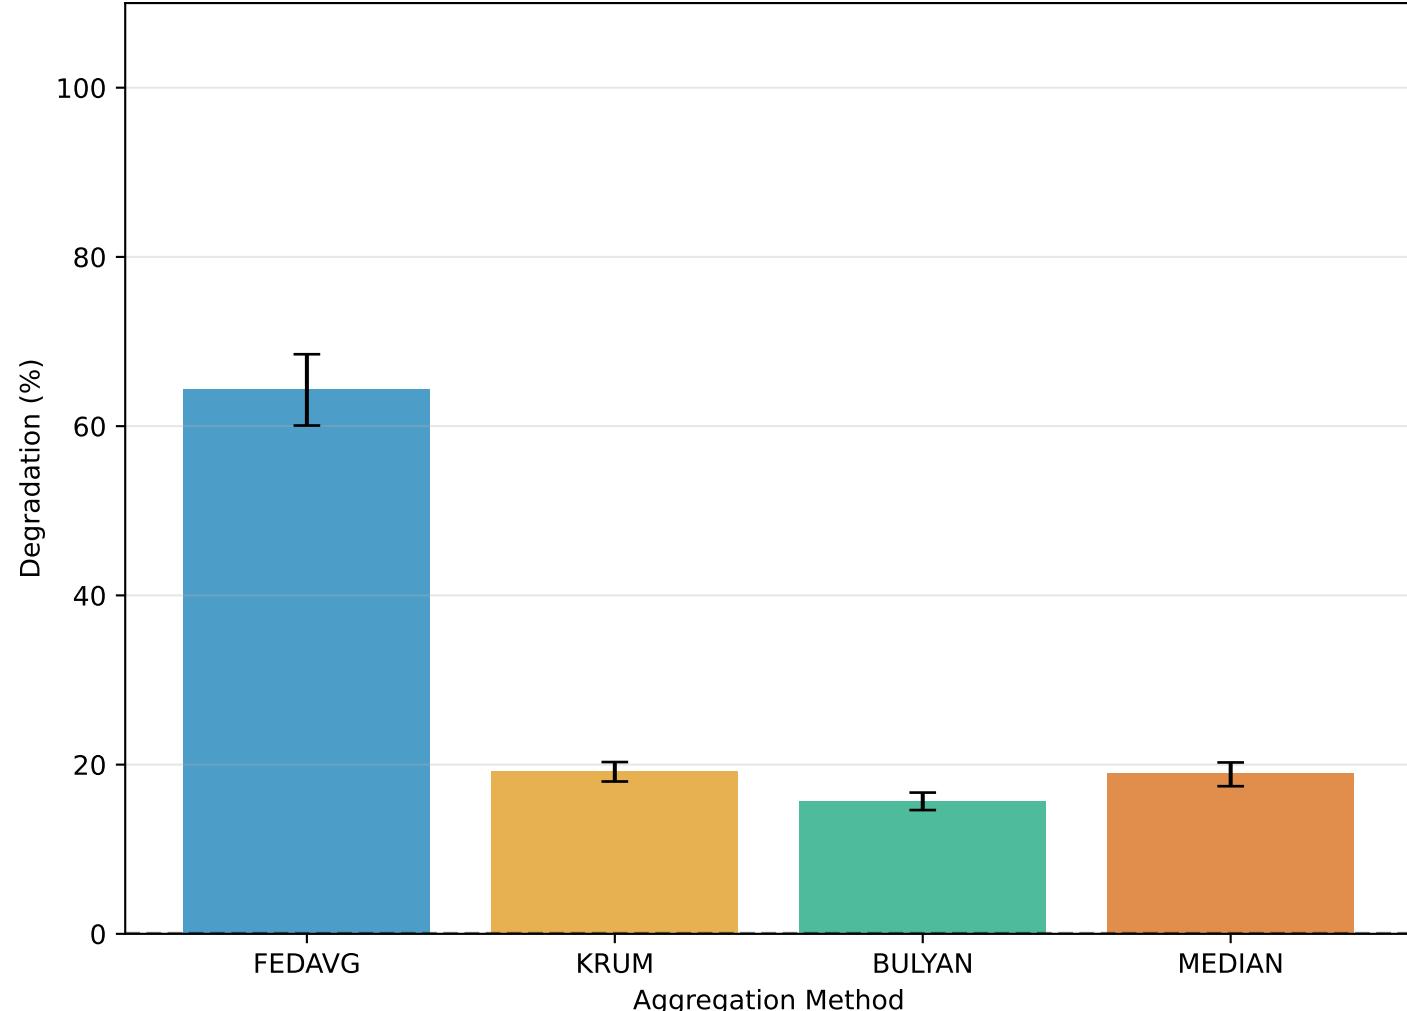
Detection Performance vs Adversary Percentage (Macro-F1, 95% CI)



Supplementary: Model Drift (L2 Distance)



Performance Degradation Under Attack (95% CI, Bounded [0,100])



Supplementary: Model Alignment (Cosine Similarity)

